

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_excel('fmcg.xlsx')
```

```
In [3]: df.head()
```

Out[3]:

	Company	Revenues (\$MM)	Profit (\$MM)	Profit as % of Revenues	Category
0	Johnson & Johnson	71890	18540	0.257894	Pharmaceuticals
1	Procter & Gamble	71726	10508	0.146502	Household & Personal Products
2	Pepsico	62789	6329	0.100798	Food
3	Pfizer	52824	7215	0.136586	Pharmaceuticals
4	Coca-Cola	41863	6527	0.155913	Beverages

```
In [4]: # Define revenue thresholds for categorical representation
small_threshold = 20000
medium_threshold = 50000

# Categorize companies based on revenue thresholds
def categorize_company_size(revenue):
    if revenue < small_threshold:
        return 'Small'
    elif revenue < medium_threshold:
        return 'Medium'
    else:
        return 'Large'

# Create new feature 'Company_Size_Categorical'
df['Company_Size_Categorical'] = df['Revenues ($MM)'].apply(categorize_company_size)
```

```
In [5]: df.head()
```

Out[5]:

	Company	Revenues (\$MM)	Profit (\$MM)	Profit as % of Revenues	Category	Company_Size_Categorical
0	Johnson & Johnson	71890	18540	0.257894	Pharmaceuticals	Large
1	Procter & Gamble	71726	10508	0.146502	Household & Personal Products	Large
2	Pepsico	62789	6329	0.100798	Food	Large
3	Pfizer	52824	7215	0.136586	Pharmaceuticals	Large
4	Coca-Cola	41863	6527	0.155913	Beverages	Medium

## Find the sum of revenue belonging to specific category

```
In [6]: # Calculate total revenue
total_revenue = df['Revenues ($MM)'].sum()

# Group by category and calculate sum of revenues
category_revenue = df.groupby('Category')['Revenues ($MM)'].sum()


# Calculate percentage of revenue for each category
percentage_revenue = (category_revenue / total_revenue) * 100

# Plot the bar plot
ax = percentage_revenue.plot(kind='bar')

# Add data labels as percentages
for i in ax.patches:
    ax.text(i.get_x() + i.get_width() / 2, i.get_height() + 0.5, f'{i.get_height():.1f}%', align='center')

# Add labels and title
plt.xlabel('Category')
plt.ylabel('Percentage of Total Revenue')
plt.title('Percentage of Total Revenue by Category')

# Show plot
plt.show()
```



## How do revenues and profits vary across different industries?

```
In [7]: # Set up the plot
plt.figure(figsize=(12, 6))

# Create side-by-side bar plots for revenues and profits by category
plt.subplot(1, 2, 1)
sns.barplot(data=df, x='Revenues ($MM)', y='Category', estimator=sum, ci=None)
plt.title('Total Revenues by Category')
plt.xlabel('Revenues ($MM)')
plt.ylabel('Category')

plt.subplot(1, 2, 2)
sns.barplot(data=df, x='Profit ($MM)', y='Category', estimator=sum, ci=None)
plt.title('Total Profits by Category')
plt.xlabel('Profit ($MM)')
plt.ylabel('')

plt.tight_layout()
plt.show()
```

C:\Users\mohan\AppData\Local\Temp\ipykernel\_16524\1699647112.py:6: FutureWarning:

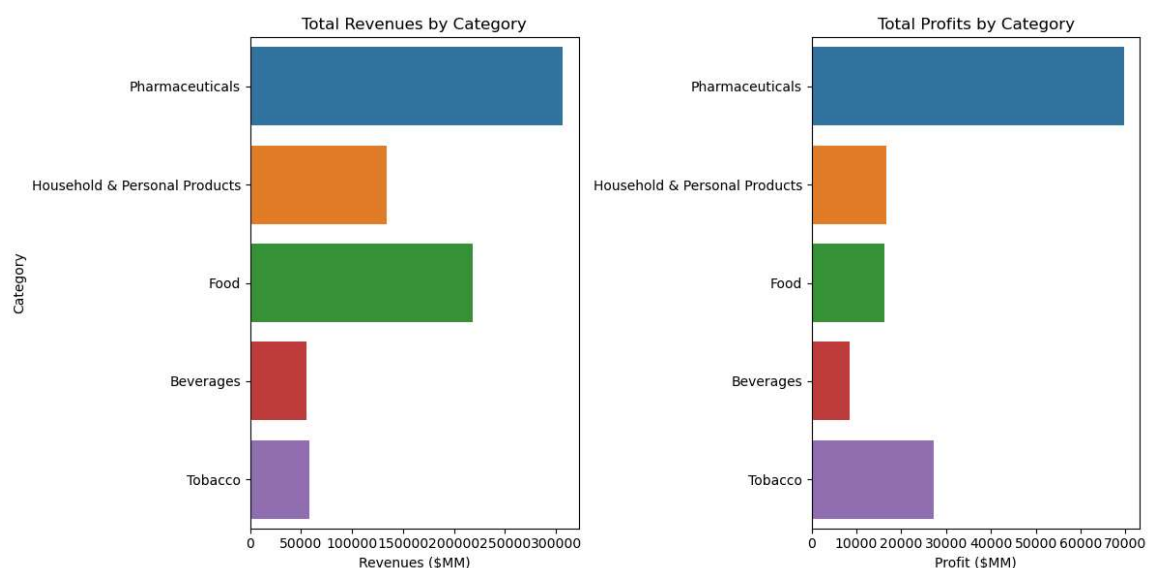
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(data=df, x='Revenues ($MM)', y='Category', estimator=sum, ci=None)
```

C:\Users\mohan\AppData\Local\Temp\ipykernel\_16524\1699647112.py:12: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(data=df, x='Profit ($MM)', y='Category', estimator=sum, ci=None)
```




## What is the distribution of profit margins within each industry?

```
In [8]: # Set up the plot
plt.figure(figsize=(10, 6))

# Create box plot for profit margins by category
sns.boxplot(data=df, x='Profit as % of Revenues', y='Category')
plt.title('Distribution of Profit Margins by Industry Category')
plt.xlabel('Profit Margin (%)')
plt.ylabel('Category')

plt.show()
```



## How do the top companies in terms of revenue compare to the top companies in terms of profit?

```
In [9]: # Sort the DataFrame by revenue and profit in descending order
df_sorted_revenue = df.sort_values(by='Revenues ($MM)', ascending=False)
df_sorted_profit = df.sort_values(by='Profit ($MM)', ascending=False)

# Get the top companies in terms of revenue and profit
top_companies_revenue = df_sorted_revenue.head(5)
top_companies_profit = df_sorted_profit.head(5)

# Merge the two sets of top companies to identify overlap
top_companies_combined = pd.merge(top_companies_revenue, top_companies_profit)

# Display the combined DataFrame
top_companies_combined
```

Out[9]:

	Company	Revenues (\$MM)_revenue	Profit (\$MM)_revenue	Profit as % of Revenues_revenue	Category_revenue	Company
0	Johnson & Johnson	71890	18540	0.257894	Pharmaceuticals	
1	Procter & Gamble	71726	10508	0.146502	Household & Personal Products	

## Is there a relationship between company size (measured by revenue) and profit margin?

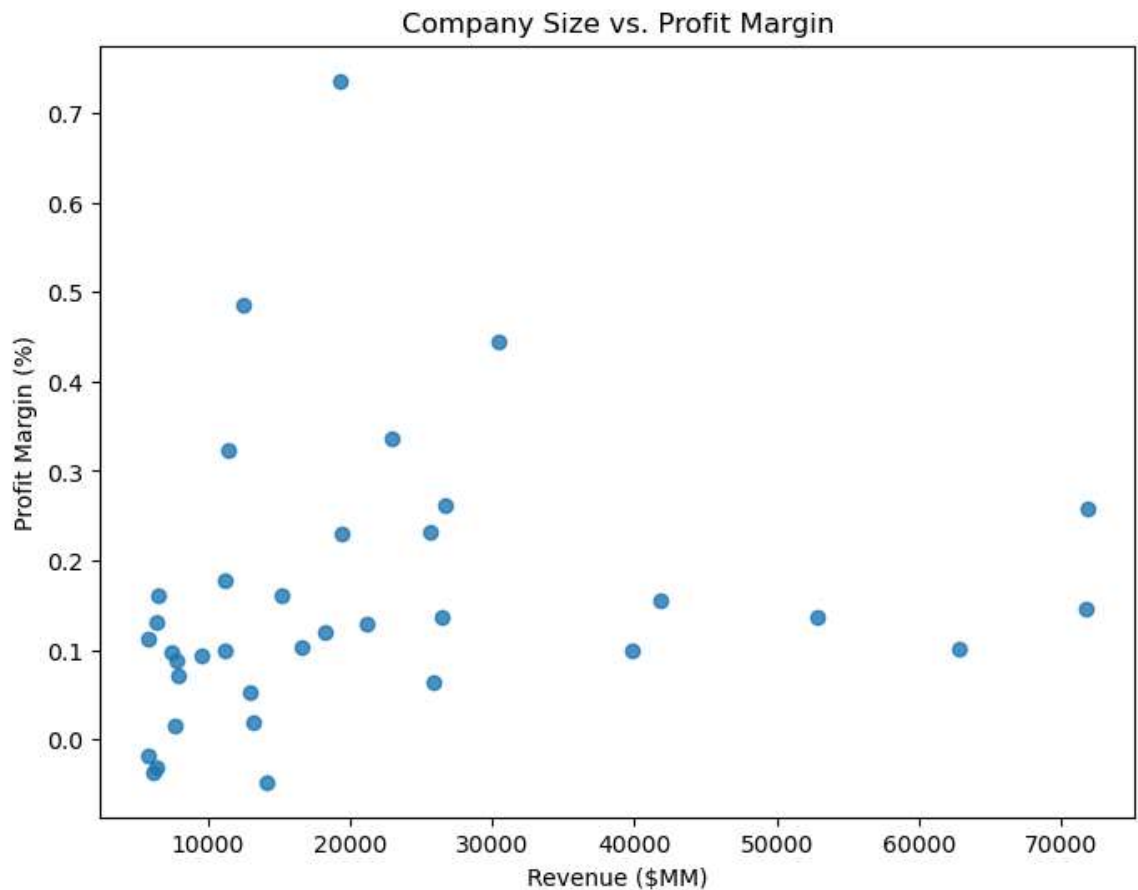
```
In [10]: # Set up the plot
plt.figure(figsize=(8, 6))

# Create scatter plot
plt.scatter(df['Revenues ($MM)'], df['Profit as % of Revenues'], alpha=0.8)
plt.title('Company Size vs. Profit Margin')
plt.xlabel('Revenue ($MM)')
plt.ylabel('Profit Margin (%)')

# Calculate correlation coefficient
correlation_coefficient = df['Revenues ($MM)'].corr(df['Profit as % of Revenues'])
print(f"Correlation Coefficient: {correlation_coefficient:.2f}")

plt.show()
```

Correlation Coefficient: 0.17



## What is the market share of each company within its respective industry?

```
In [11]: # Calculate total revenue for each industry category
total_revenue_by_category = df.groupby('Category')['Revenues ($MM)'].sum()

# Calculate market share for each company within its industry category
df['Market Share'] = df.apply(lambda row: row['Revenues ($MM)'] / total_revenue_by_category[row['Category']], axis=1)

# Display the DataFrame with market share
df[['Company', 'Category', 'Revenues ($MM)', 'Market Share']]
```




Out[11]:

	Company	Category	Revenues (\$MM)	Market Share
0	Johnson & Johnson	Pharmaceuticals	71890	0.234271
1	Procter & Gamble	Household & Personal Products	71726	0.533668
2	Pepsico	Food	62789	0.287018
3	Pfizer	Pharmaceuticals	52824	0.172140
4	Coca-Cola	Beverages	41863	0.763213
5	Merck	Pharmaceuticals	39807	0.129721
6	Gilead Sciences	Pharmaceuticals	30390	0.099033
7	Philip Morris International	Tobacco	26685	0.455959
8	Kraft Heinz	Food	26487	0.121076
9	Mondelez International	Food	25923	0.118498
10	Abbvie	Pharmaceuticals	25638	0.083548
11	Amgen	Pharmaceuticals	22991	0.074922
12	Eli Lilly	Pharmaceuticals	21222	0.069157
13	Bristol-Myers Squibb	Pharmaceuticals	19427	0.063308
14	Altria Group	Tobacco	19337	0.330406
15	Kimberly-Clark	Household & Personal Products	18202	0.135430
16	General Mills	Food	16563	0.075712
17	Colgate-Palmolive	Household & Personal Products	15195	0.113056
18	Conagra Brands	Food	14134	0.064609
19	Land O'Lakes	Food	13233	0.060490
20	Kellogg	Food	13014	0.059489
21	Reynolds American	Tobacco	12503	0.213635
22	Biogen	Pharmaceuticals	11449	0.037309
23	Estee Lauder	Household & Personal Products	11262	0.083793
24	Celgene	Pharmaceuticals	11229	0.036592
25	Hormel Foods	Food	9523	0.043531
26	Campbell Soup	Food	7961	0.036391
27	J. M. Smucker	Food	7811	0.035705
28	Dean Foods	Food	7710	0.035244
29	Hershey	Food	7440	0.034009
30	Constellation Brands	Beverages	6548	0.119378
31	Dr. Pepper Snapple Group	Beverages	6440	0.117409
32	HRG Group	Household & Personal Products	6403	0.047641
33	Treehouse Foods	Food	6175	0.028227
34	Avon Products	Household & Personal Products	5853	0.043548
35	Clorox	Household & Personal Products	5761	0.042864

```
In [12]: df.to_csv('fmcg_main.csv')
```

```
In [13]: df=pd.read_csv('fmcg_main.csv')
# Plotting the relationship between company size and profitability
sns.boxplot(x='Company_Size_Categorical', y='Profit as % of Revenues', data=
plt.title('Profitability by Company Size')
plt.xlabel('Company Size')
plt.ylabel('Profit as % of Revenues')
plt.show()
```



```
In [14]: import pandas as pd
from scipy.stats import f_oneway

# Extract profitability data for each company size category
small = df[df['Company_Size_Categorical'] == 'Small']['Profit as % of Revenue']
medium = df[df['Company_Size_Categorical'] == 'Medium']['Profit as % of Revenue']
large = df[df['Company_Size_Categorical'] == 'Large']['Profit as % of Revenue']

# Perform ANOVA test
statistic, p_value = f_oneway(small, medium, large)

# Set significance level
alpha = 0.05

# Print results
print("ANOVA Test Results:")
print("Test Statistic:", statistic)
print("p-value:", p_value)

# Interpret results
if p_value < alpha:
    print("There are significant differences in profitability among companies of different sizes.")
else:
    print("There are no significant differences in profitability among companies of different sizes.")
```

ANOVA Test Results:  
Test Statistic: 0.62917886057582  
p-value: 0.5393020854290076  
There are no significant differences in profitability among companies of different sizes.

```
In [15]: # Calculate average profitability for each category
average_profitability = df.groupby('Category')['Profit as % of Revenues'].mean()

# Identify the most profitable category
most_profitable_category = average_profitability.idxmax()
average_profitability_value = average_profitability.max()

print("Most Profitable Category:", most_profitable_category)
print("Average Profitability:", average_profitability_value)
```

Most Profitable Category: Tobacco  
Average Profitability: 0.4943889257543861

In [16]:

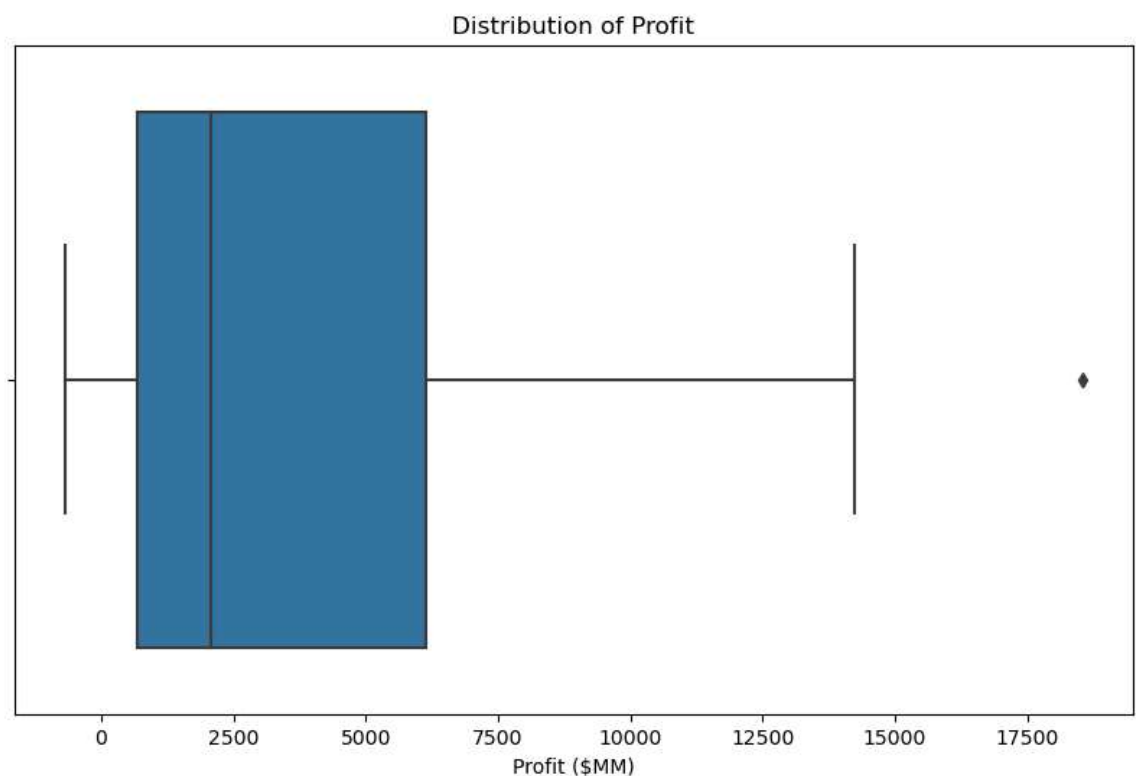
```
# Visualize distribution of profit using a box plot
plt.figure(figsize=(10, 6))
sns.boxplot(x='Profit ($MM)', data=df)
plt.title('Distribution of Profit')
plt.show()


# Visualize distribution of revenue using a box plot
plt.figure(figsize=(10, 6))
sns.boxplot(x='Revenues ($MM)', data=df)
plt.title('Distribution of Revenue')
plt.show()

# Calculate z-scores for profit and revenue
df['Profit Z-Score'] = (df['Profit ($MM)'] - df['Profit ($MM)'].mean()) / df['Profit ($MM)'].std()
df['Revenue Z-Score'] = (df['Revenues ($MM)'] - df['Revenues ($MM)'].mean()) / df['Revenues ($MM)'].std()

# Identify outliers based on z-scores (threshold of |z-score| > 3)
profit_outliers = df[abs(df['Profit Z-Score']) > 3]
revenue_outliers = df[abs(df['Revenue Z-Score']) > 3]

print("Profit Outliers:")
print(profit_outliers)
print("\nRevenue Outliers:")
print(revenue_outliers)
```





Profit Outliers:

Unnamed: 0	Company	Revenues (\$MM)	Profit (\$MM)	\
0	Johnson & Johnson	71890	18540	

Profit as % of Revenues	Category	Company_Size_Categorical	\
0.257894	Pharmaceuticals	Large	

Market Share	Profit Z-Score	Revenue Z-Score
0.234271	3.242016	2.763337

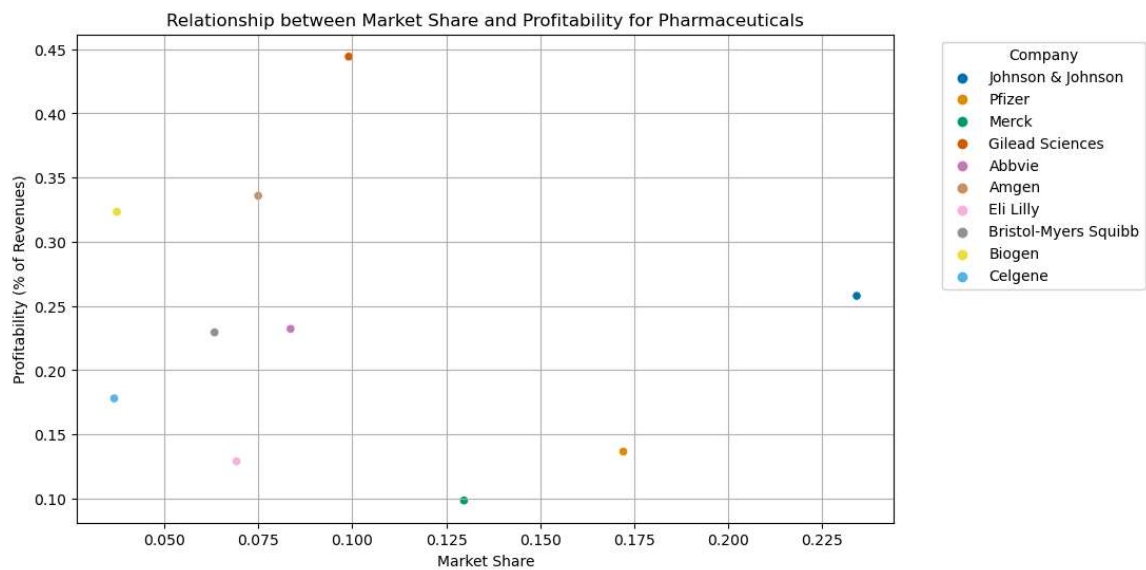
Revenue Outliers:

Empty DataFrame  
Columns: [Unnamed: 0, Company, Revenues (\$MM), Profit (\$MM), Profit as % of Revenues, Category, Company\_Size\_Categorical, Market Share, Profit Z-Score, Revenue Z-Score]  
Index: []

```
In [17]: # Select a specific category for analysis (e.g., Pharmaceuticals)
category = 'Pharmaceuticals'

# Filter data for the selected category
category_data = df[df['Category'] == category]

# Visualize the relationship between market share and profitability using a
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Market Share', y='Profit as % of Revenues', data=category_data)
plt.title('Relationship between Market Share and Profitability for {}'.format(category))
plt.xlabel('Market Share')
plt.ylabel('Profitability (% of Revenues)')
plt.legend(title='Company', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(True)
plt.show()
```



In [ ]: